

# A Simple Index for Wikipedia Page Retrieval

Murari Mani



# Objectives

- Build a reverse index for quick lookup for wikipedia
- Rank documents according to query and document contents
- Retrieve best document matching query

# Data Source

- Wikimedia data dump of current articles  
pages-articles-multistream.xml.bz2
  - Current revisions only, no talk or user pages
- <http://itorrents.org/torrent/3A8C87DE09C85193CFBCB10DC64B7A64C2CEE7FC.torrent> Size: ~15 GB compressed
- Tested by downloading and extracting the simple english wikipedia which is 179 MB uncompressed
  - simplewiki-20170820-pages-meta-current.xml.bz2 (179.8 MiB) on itorrents.org
- Parser : <https://github.com/attardi/wikiextractor>

# Structure of Parsed Data

Split into xml pages of 100MB each (configurable) with individual pages delineated by <doc> and </doc> tags

```
<doc id="1" url="https://simple.wikipedia.org/wiki?curid=1" title="April">
```

April

April is the fourth month of the year, and comes between March and May. It is one of four months to have 30 days.

...

```
</doc>
```

# Strategy for Storage & Document Retrieval

**Step 1:** Read data into HDFS

**Natural choice since the data is unstructured and we will be performing a bunch of map-reduce tasks**

**Step 2:** Compute the tf-idf\* score to get the most relevant meta-document i.e wiki\_00

**Step 3:** Within wiki\_00 compute the tf-idf\* for individual html pages to get the best matching webpage

\*<https://nlp.stanford.edu/IR-book/pdf/06vect.pdf>

# Strategy for Storage & Document Retrieval

**Step 1:** Read data into HDFS

**Natural choice since the data is unstructured and we will be performing a bunch of map-reduce tasks**

**Step 2:** Compute the tf-idf score to get the most relevant meta-document i.e wiki\_00

**Step 3:** Within wiki\_00 compute the tf-idf for individual html pages to get the best matching webpage

# Shortcomings

1. Performs well for uncommon queries such as “blah”
2. Does not work so well for common queries such as “president”
3. Reason: All fields of the document are equally weighted
4. We should assign more weight to the title/summary so that the score is adjusted accordingly
5. **Errors when I try to run the large data set on the SandBox**

# Future Extensions

1. Support more complex queries
2. Support zone based weighting
3. Improve the ranking by incorporating more sophisticated ML-based ranking
4. Deploy on a real HDFS cluster

Demo at: <https://youtu.be/leZVQ9L9j-g>





Thanks for listening :) !!